Correspondence

Listen to Gen Z on AI in education

In his World View, Mim Rahimi addresses a growing distrust towards artificial intelligence in Generation Z university students and researchers (see *Nature* **654**, 299; 2026).

He feels that universities should respond by teaching AI literacy. Students should be taught how to verify AI claims and when they are susceptible to hallucination. It is, for Rahimi, "not enough to offer students free AI subscriptions and leave them to work out the implications of its use".

We do not deny that such literacy is important, but it does nothing to assuage the concerns among members of Gen Z — usually defined as people born between 1997 and 2012 — about AI in education. Rahimi runs the risk of unfairly painting these legitimate worries as stemming from a misguided technophobia. He says that without AI literacy, Gen Z students are engaged in a "technological shift that they are not ready to treat as neutral". But this assumes that when they are educated about AI, students will drop their scepticism and treat AI and its role in their education 'neutrally'. This is not obvious.

It is hard to acquire the skills to verify claims when AI use is permitted. AI hallucinations are not predictable. And if students are allowed to use AI liberally when crafting literature reviews, they are unlikely to learn how to situate their findings in a broader scientific context.

In responding to Gen Z AI scepticism, we must be open to the possibility that the concerns are well founded.

Ian Robertson, **Vincent C. Müller**, **Eleonora Catena** Friedrich-Alexander University of Erlangen–Nuremberg, Erlangen, Germany.
ian.george.robertson@fau.de

AI can cause harm: safeguards must catch up

In a World View last year, I argued that emotionally responsive artificial intelligence needs mandatory safeguards (see *Nature* **643**, 9; 2025). Several events have honed this point.

Last month, AI developer Anthropic in San Francisco, California, took its most capable AI models offline within days of their release, citing a US government directive to suspend access by any foreign national owing to cybersecurity risks. The export controls have since been lifted.

Families are pursuing lawsuits over deaths linked to chatbot use. In controlled studies, AI agents carried out instructions to cause harm in experimental settings even while judging their own behaviour as immoral (H. Dery *et al.* Preprint at PsyArXiv https://doi.org/rb9x (2026); Z. Ben-Zion *et al.* Preprint at https://doi.org/rcbb; 2026).

The responses to these events have been strikingly different. A speculative risk in the case of the Anthropic resulted in a speedy global takedown. Yet, cases about deaths might spend years in litigation before any action can be taken. Furthermore, a California state law mandating safeguards for chatbots might be challenged by a US executive order that aims to reduce AI regulation.

There is no question that AI can recognize harm. The discussion must move to whether AI can stop when it notices that its behaviour is dangerous, whether we will require it to and which harms are treated as emergencies.

Ziv Ben-Zion University of Haifa, Haifa, Israel.
zbenzion@univ.haifa.ac.il

The author declares competing interests; see go.nature.com/3qrksgu for details.

How long can humans live? Up to 125 years

In your interview with longevity researcher Saul Newman, he states that our paper argues that there is a "hard limit to human survival" but is based on "basic mathematical errors" (see *Nature* **654**, 29–31; 2026). Both statements are incorrect.

Our paper gives a nuanced appraisal of the limit to human lifespan, noting that the oldest person in any year is likely to be around 115 years old, with a 95% confidence interval of 113.1–116.7 years old (X. Dong *et al. Nature* **538**, 257–259; 2016). Our finding has held up: the current oldest living person is 116 years old . We also calculated a firmer limit of 125 years, which has a less than 1 in 10,000 chance of being broken. The human maximum lifespan has not risen since the 1990s. Conflating these findings as a "hard limit" misrepresents them.

We have previously debunked assertions of mathematical errors (X. Dong *et al. Nature* **546**, E12; 2017). Out of thousands of verified supercentenarians — people aged 110 or older — only one has lived past 120. But in any case, even if few or no people have lived past the age of 110, that only lends support to the idea that there could be a limit to human lifespan.

Brandon Milholland IQVIA, Durham, North Carolina, USA.

Xiao Dong University of Minnesota, Minneapolis, Minnesota, USA.

Jan Vijg Albert Einstein College of Medicine, New York City New York, USA.
jan.vijg@einsteinmed.edu

J.V. declares competing interests; see go.nature.com/4guyh1j for details.

HOW TO SUBMIT CORRESPONDENCE

Please consult the full author guidelines and policies at go.nature.com/cmchno before e-mailing your submission to correspondence@nature.com. It should not be sent as an attachment.

The following cannot be considered: technical comments on peer-reviewed research papers; responses to articles published in journals other than *Nature*; contributions that present primary research data; and submissions that do not comply with the section's strict length and style requirements. Submissions must not be under consideration elsewhere.

Please make clear in your proposed text its pertinence to *Nature*'s global readership. Ensure that you include a link to any article under discussion and references or links for fact-checking of all statements, in addition to the limited number of references necessary for publication.